

# Estimating a Poisson Regression Model with R

Rémi Piatek

May 6, 2019

# The R Programming Language

- ▶ R: Popular **open-source programming language** for statistical analysis.
- ▶ Widely used in statistics and econometrics.
- ▶ **User-friendly and powerful IDE** for R: RStudio.
- ▶ Basic functionalities of R can be extended by **packages**.
- ▶ Large number of packages available on the Comprehensive R Archive Network.
- ▶ **Goal of this presentation:** Illustrate how to use R for the estimation of a Poisson regression model.

# Count Data Models

- ▶ **Count data** models are used to explain dependent variables that are natural numbers, i.e., positive integers such that  $y_i \in \mathbb{N}$ , where  $\mathbb{N} = \{0, 1, 2, \dots\}$ .
- ▶ Count data models are frequently used in economics to study **countable events**: Number of years of education, number of patent applications filed by companies, number of doctor visits, number of crimes committed in a given city, etc.
- ▶ The **Poisson model** is a popular count data model.

# Poisson Regression Model

- ▶ Given a parameter  $\lambda_i > 0$ , the **Poisson model** assumes that the probability of observing  $Y_i = y_i$ , where  $y_i \in \mathbb{N}$ , is equal to:

$$\text{Prob}(Y_i = y_i \mid \lambda_i) = \frac{\lambda_i^{y_i} \exp\{-\lambda_i\}}{y_i!},$$

for  $i = 1, \dots, N$ .

- ▶ The mean and the variance of  $Y_i$  are equal to the parameter  $\lambda_i$ :

$$E(Y_i \mid \lambda_i) = V(Y_i \mid \lambda_i) = \lambda_i,$$

implying *equi-dispersion* of the data.

- ▶ To control for **observed characteristics**, the parameter  $\lambda_i$  can be parametrized as follows (implying  $\lambda_i > 0$ ):

$$E(Y_i \mid X_i, \beta) \equiv \lambda_i = \exp\{X_i' \beta\},$$

where  $X_i$  is a vector containing the covariates.

## Simulating Data

- ▶ R function simulating data from Poisson regression model:

```
simul_poisson <- function(n, beta) {  
  k <- length(beta)           # number of covariates  
  x <- replicate(k - 1, rnorm(n)) # simulate covariates  
  x <- cbind(1, x)            # for intercept term  
  lambda <- exp(x %*% beta)    # individual means  
  y <- rpois(n, lambda)        # simulate count  
  return(data.frame(y, x))     # return variables  
}
```

- ▶ Using function to generate data:

```
set.seed(123)  
nobs <- 1000  
beta <- c(-.5, .4, -.7)  
data <- simul_poisson(nobs, beta)
```

# Data Description

## ► Descriptive statistics:

```
# extract variables of interest from data set
y <- data[, 1]
x <- as.matrix(data[, 2:4])


# descriptive statistics
library(psych)
describe(data)
```

```
##      vars      n mean   sd median trimmed  mad   min   max ra
## y         1 1000 0.76 1.08   0.00     0.54 0.00   0.00 8.00 8
## X1        2 1000 1.00 0.00   1.00     1.00 0.00   1.00 1.00 0
## X2        3 1000 0.02 0.99   0.01     0.01 0.96 -2.81 3.24 6
## X3        4 1000 0.04 1.01   0.05     0.05 1.05 -3.05 3.39 6
##              se
## y    0.03
## X1  0.00
## X2  0.03
```

# Data Description

- Histogram of count variable:



```
barplot(table(y))
```



# Data Description

- Relationship between count variable and covariates:

```
par(mfrow = c(1, 2))  
plot(y, x[, 2])  
plot(y, x[, 3])
```





# Likelihood Function and ML Estimator

- ▶ Individual contribution to the likelihood function:

$$L_i(\beta; y_i, x_i) = \frac{\exp\{y_i x_i \beta\} \exp\{-\exp\{x_i \beta\}\}}{y_i!}$$

- ▶ Individual log-Likelihood function:

$$\ell_i(\beta; y_i, x_i) = \log L_i(\beta; y_i, x_i) = y_i x_i \beta - \exp\{x_i \beta\} - \log(y_i!)$$

- ▶ Maximum Likelihood Estimator:

$$\hat{\beta}_{\text{MLE}} = \arg \max_{\beta} \sum_{i=1}^N \ell(\beta; y, X)$$

- ▶ Optimization (using *minimization* of objective function):

$$\hat{\beta}_{\text{MLE}} = \arg \min_{\beta} Q(\beta; y, X) \quad Q(\beta; y, X) = -\frac{1}{N} \sum_{i=1}^N \ell_i(\beta; y_i, x_i)$$

## Coding the Objective Function

```
# Objective function of Poisson regression model
obj_poisson <- function(beta, y, x) {
  lambda <- x %*% beta
  llik <- y*lambda - exp(lambda) - lfactorial(y)
  return(-mean(llik))
}
```

```
# Evaluating objective function
beta0 <- c(1, 2, 3)
obj_poisson(beta0, y, x)
```

```
## [1] 1757.113
```

# Maximizing the Objective Function

- ▶ Set starting values:

```
beta0 <- rep(0, length(beta))
```

- ▶ Optimize using quasi-Newton method (BFGS algorithm):

```
opt <- optim(beta0, obj_poisson, method = "BFGS",  
             y = y, x = x)
```

- ▶ Show results:

```
cat("ML estimates:", opt$par,  
    "\nObjective function:", opt$value, "\n")
```

```
## ML estimates: -0.5740286 0.3921569 -0.7231029  
## Objective function: 0.9998689
```

## Comparing Results to Built-in Function

```
opt_glm <- glm(y ~ 0 + x, family = poisson)
summary(opt_glm)
```

```
##
```

```
## Call:
```

```
## glm(formula = y ~ 0 + x, family = poisson)
```

```
##
```

```
## Deviance Residuals:
```

```
##      Min        1Q      Median        3Q        Max
## -2.1707  -0.9185  -0.5915    0.5133    3.6768
```

```
##
```

```
## Coefficients:
```

```
##      Estimate Std. Error z value Pr(>|z|)
## xX1 -0.57403    0.04602  -12.47  <2e-16 ***
## xX2  0.39216    0.03690   10.63  <2e-16 ***
## xX3 -0.72310    0.03662  -19.74  <2e-16 ***
```

```
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1
```

## Comparing Results to Built-in Function

- Collect results from the two approaches to compare them:

```
res <- cbind("True" = beta, "MLE" = opt$par,  
            "GLM" = opt_glm$coefficients)  
row.names(res) <- c("constant", "x1", "x2")  
res
```

##	True	MLE	GLM
## constant	-0.5	-0.5740286	-0.5740289
## x1	0.4	0.3921569	0.3921561
## x2	-0.7	-0.7231029	-0.7231039

- **Question:** Our results (MLE) are virtually the same as those obtained with the built-in function GLM, but not identical. Where do the small differences come from?

# Empirical Illustration

- ▶ Goal: Investigate the determinants of fertility.
- ▶ Poisson regression model used to estimate the relationship between explanatory variables and count outcome variable.
- ▶ Both our estimator coded from scratch and R built-in function will be used.

# Data

- ▶ Source: Botswana's 1988 Demographic and Health Survey.
- ▶ Data set borrowed from Wooldridge:

```
library(wooldridge)  
data(fertil2)
```

- ▶ Outcome variable: Total number of living children:

```
y_lab <- "children"
```

- ▶ Explanatory variables: Education, age, marital status, living in urban area, having electricity/TV at home:

```
x_lab <- c("educ", "age", "agesq", "evermarr", "urban",  
          "electric", "tv")
```

## Loading data

- ▶ Selecting variables and removing missing values:

```
data <- fertil2[, c(y_lab, x_lab)]  
data <- na.omit(data)
```

- ▶ Show first 6 observations on first 8 variables:

```
head(data[, 1:8], n = 6)
```

##	children	educ	age	agesq	evermarr	urban	electric	tv
## 1	0	12	24	576	0	1	1	1
## 2	3	13	32	1024	1	1	1	1
## 3	1	5	30	900	1	1	1	0
## 4	2	4	42	1764	1	1	1	1
## 5	2	11	43	1849	1	1	1	1
## 6	1	7	36	1296	1	1	1	0



# Descriptive Statistics



```
library(psych)
```

```
describe(data)
```

##		vars	n	mean	sd	median	trimmed	mad
##	children	1	4358	2.27	2.22	2	1.95	2.97
##	educ	2	4358	5.86	3.93	7	5.79	4.45
##	age	3	4358	27.40	8.69	26	26.71	8.90
##	agesq	4	4358	826.42	526.96	676	752.87	467.02
##	evermarr	5	4358	0.48	0.50	0	0.47	0.00
##	urban	6	4358	0.52	0.50	1	0.52	0.00
##	electric	7	4358	0.14	0.35	0	0.05	0.00
##	tv	8	4358	0.09	0.29	0	0.00	0.00
##		skew	kurtosis	se				
##	children	1.07	0.75	0.03				
##	educ	-0.03	-0.50	0.06				
##	age	0.59	-0.49	0.13				
##	agesq	1.10	0.52	7.98				
##	evermarr	0.09	-1.99	0.01				

# Plot

```
attach(data)
par(mfrow = c(1, 2))
blue_transp <- adjustcolor("blue", alpha.f = 0.1)
plot(age, children, pch = 19, col = blue_transp)
plot(educ, children, pch = 19, col = blue_transp)
```



# MLE of the Poisson Model

- ▶ Maximum likelihood function using built-in function `glm()`:

```
mle <- glm(children ~ educ + age + agesq + evermarr +  
            urban + electric + tv,  
            family = "poisson", data = data)
```

- ▶ Maximum likelihood function using our own function:

```
y <- data[, y_lab]  
x <- as.matrix(data[, x_lab])  
x <- cbind(1, x)           # for intercept term  
beta0 <- rep(0, ncol(x))  # starting values  
opt <- optim(beta0, obj_poisson, method = "BFGS",  
             y = y, x = x)
```

# MLE of the Poisson Model

- ▶ Results different from `glm()`?
- ▶ Optimization algorithms are iterative methods that rely on different criteria to determine if/when the optimum has been reached.
- ▶ For example: Change in the objective function, change in the parameter values, change in the gradient, step size, etc.
- ▶ *[More in Advanced Microeconometrics course]*.
- ▶ Try to adjust tuning parameters, for example add `control = list(ndeps = rep(1e-8, ncol(x)))` to `optim()` to change step size of gradient approximation.

## Summarizing the Empirical Results

```
summary(mle)
```

```
##
```

```
## Call:
```

```
## glm(formula = children ~ educ + age + agesq + evermarr -
```

```
##      electric + tv, family = "poisson", data = data)
```

```
##
```

```
## Deviance Residuals:
```

```
##      Min        1Q      Median        3Q        Max
```

```
## -3.5620  -0.8116  -0.1091   0.5439   2.8893
```

```
##
```

```
## Coefficients:
```

```
##              Estimate Std. Error z value Pr(>|z|)
```

```
## (Intercept) -5.3748294  0.1628671 -33.001  < 2e-16 ***
```

```
## educ        -0.0216645  0.0029131  -7.437 1.03e-13 ***
```


```
## age         0.3373308  0.0099365  33.949 < 2e-16 ***
```

```
## agesq      -0.0041158  0.0001453 -28.331 < 2e-16 ***
```

```
## evermarr    0.3147510  0.0244473  12.875 < 2e-16 ***
```

# Fitted Values

```
plot(density(mle$fitted.values),  
     main = "Density of fitted mean values")
```



## Formatting the results

```
library(xtable)
```

```
xtable(mle)
```

```
## % latex table generated in R 3.6.0 by xtable 1.8-4 packa
```

```
## % Mon May 6 11:06:08 2019
```

```
## \begin{table}[ht]
```

```
## \centering
```

```
## \begin{tabular}{rrrrr}
```

```
## \hline
```

```
## & Estimate & Std. Error & z value & Pr(>|z|) & \\\
```

```
## \hline
```

```
## (Intercept) & -5.3748 & 0.1629 & -33.00 & 0.0000 \\\
```

```
## educ & -0.0217 & 0.0029 & -7.44 & 0.0000 \\\
```

```
## age & 0.3373 & 0.0099 & 33.95 & 0.0000 \\\
```

```
## agesq & -0.0041 & 0.0001 & -28.33 & 0.0000 \\\
```

```
## evermarr & 0.3148 & 0.0244 & 12.87 & 0.0000 \\\
```

```
## urban & -0.0861 & 0.0216 & -3.98 & 0.0001 \\\
```

```
## electric & -0.1205 & 0.0388 & -3.10 & 0.0019 \\\
```